

Module task cases

hidden targets are scorer-only; the agent receives only the visible input

A Recovery



visible input C^{in}



recover



complete module

B Explanation



visible input $C^{\text{in}} = C$



explain, do not alter

mechanistic claim

evidence handles

C Refinement



visible input C^{in}



remove distractors



refined module

D No mechanism



visible pseudo-module



calibrate

split groups

abstain

never fabricate one module

Corpus difficulty

harder distractors lie closer to the target

